

```
from google.colab import drive
drive.mount('/content/drive')
```

```
base_dir = '/content/drive/My Drive/Senior Year/Applied Data Science'
metadata_path = f'{base_dir}/sample_reduced_features.csv'
image_dir = f'{base_dir}/Data/downsampled_data/'
```

↻ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)

```
import pandas as pd
import numpy as np
import tensorflow as tf
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from tensorflow.keras.layers import Input, Dense, Flatten, Concatenate
from tensorflow.keras.models import Model
```

```
metadata_path = '/content/drive/My Drive/Senior Year/Applied Data Science/sample_reduced_features.csv'
metadata = pd.read_csv(metadata_path)
```

```
image_dir = '/content/drive/My Drive/Senior Year/Applied Data Science/Data/downsampled_data/'
```

```
metadata['Image_Path'] = metadata['Simplified_FileName'] + '_median_aggregated.tiff'
image_paths = image_dir + metadata['Image_Path']
```

```
meta_features = ['area', 'mean_intensity', 'bbox-0', 'bbox-1', 'bbox-2', 'bbox-3', 'eccentricity', 'solidity']
X_meta = metadata[meta_features]
y_labels = metadata['Metadata_pert_iname']
```

```
label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y_labels)
```

```
scaler = StandardScaler()
X_meta_scaled = scaler.fit_transform(X_meta)
```

```
X_meta_train, X_meta_val, y_train, y_val = train_test_split(X_meta_scaled, y_encoded, test_size=0.2, random_state=42)
image_paths_train, image_paths_val = train_test_split(image_paths, test_size=0.2, random_state=42)
```

```
def load_image(path):
    img = tf.keras.utils.load_img(path.numpy().decode('utf-8'), target_size=(64, 64))
    img = tf.keras.utils.img_to_array(img) / 255.0
    return img
```

```
def hybrid_data_tf_generator(image_paths, X_meta, labels, batch_size=32):
    image_paths = tf.convert_to_tensor(image_paths.values, dtype=tf.string)
    dataset = tf.data.Dataset.from_tensor_slices((image_paths, X_meta, labels))
```

```
    def process_data(img_path, meta_data, label):
        img = tf.py_function(func=load_image, inp=[img_path], Tout=tf.float32)
        img.set_shape((64, 64, 3))
        return {"image_input": img, "meta_input": meta_data}, label
```

```
    dataset = dataset.map(process_data, num_parallel_calls=tf.data.AUTOTUNE)
    dataset = dataset.shuffle(buffer_size=1000).batch(batch_size).prefetch(tf.data.AUTOTUNE)
    return dataset
```

```
image_input = Input(shape=(64, 64, 3), name="image_input")
x = tf.keras.layers.Conv2D(32, (3, 3), activation='relu')(image_input)
x = Flatten()(x)
```

```
meta_input = Input(shape=(X_meta_train.shape[1],), name="meta_input")
meta_features = Dense(32, activation='relu')(meta_input)
```

```
combined = Concatenate()([x, meta_features])
output = Dense(len(label_encoder.classes_), activation='softmax')(combined)
```

```
hybrid_model = Model(inputs={"image_input": image_input, "meta_input": meta_input}, outputs=output)
hybrid_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

```
train_dataset = hybrid_data_tf_generator(image_paths_train, X_meta_train, y_train, batch_size=32)
val_dataset = hybrid_data_tf_generator(image_paths_val, X_meta_val, y_val, batch_size=32)
```

```
history = hybrid_model.fit(
    train_dataset,
    validation_data=val_dataset
```

```

validation_data=val_dataset,
epochs=30,
callbacks=[
    tf.keras.callbacks.EarlyStopping(patience=5, restore_best_weights=True),
    tf.keras.callbacks.ReduceLROnPlateau(factor=0.2, patience=3, min_lr=1e-5)
],
verbose=1
)

```

```

Epoch 1/30
363/363 ————— 73s 181ms/step - accuracy: 0.4598 - loss: 3.4246 - val_accuracy: 0.9697 - val_loss: 0.2843 - le
Epoch 2/30
363/363 ————— 56s 143ms/step - accuracy: 0.9963 - loss: 0.0352 - val_accuracy: 0.9752 - val_loss: 0.2651 - le
Epoch 3/30
363/363 ————— 56s 143ms/step - accuracy: 0.9992 - loss: 0.0118 - val_accuracy: 0.9759 - val_loss: 0.2795 - le
Epoch 4/30
363/363 ————— 56s 145ms/step - accuracy: 0.9995 - loss: 0.0079 - val_accuracy: 0.9772 - val_loss: 0.2933 - le
Epoch 5/30
363/363 ————— 57s 145ms/step - accuracy: 0.9997 - loss: 0.0054 - val_accuracy: 0.9772 - val_loss: 0.2950 - le
Epoch 6/30
363/363 ————— 56s 143ms/step - accuracy: 0.9995 - loss: 0.0078 - val_accuracy: 0.9772 - val_loss: 0.2967 - le
Epoch 7/30
363/363 ————— 55s 142ms/step - accuracy: 0.9995 - loss: 0.0082 - val_accuracy: 0.9772 - val_loss: 0.2977 - le

```

```

import os
import numpy as np
import tensorflow as tf
from tensorflow.keras.utils import load_img, img_to_array

example_image_dir = '/content/drive/My Drive/Senior Year/Applied Data Science/Data/example_data/'

example_image_paths = [os.path.join(example_image_dir, fname) for fname in os.listdir(example_image_dir) if fname.endswith('.tif')]

def preprocess_images(paths):
    images = []
    for path in paths:
        img = load_img(path, target_size=(64, 64))
        img_array = img_to_array(img) / 255.0
        images.append(img_array)
    return np.array(images)

X_image_example = preprocess_images(example_image_paths)

dummy_meta = np.zeros((len(X_image_example), X_meta_train.shape[1]))

predictions = hybrid_model.predict({"image_input": X_image_example, "meta_input": dummy_meta})

predicted_drugs = label_encoder.inverse_transform(np.argmax(predictions, axis=1))

for i, path in enumerate(example_image_paths):
    print(f"Image: {os.path.basename(path)}, Predicted Drug: {predicted_drugs[i]}")

```

```

1/1 ————— 1s 856ms/step
Image: r04c08f05p01-compound-FK866.tiff, Predicted Drug: CC-401
Image: r04c14f05p01-compound-DMSO.tiff, Predicted Drug: CC-401
Image: r06c10f05p01-compound-quinidine.tiff, Predicted Drug: CC-401
Image: r12c09f05p01-compound-FK866.tiff, Predicted Drug: CC-401
Image: r13c02f05p01-compound-LY2109761.tiff, Predicted Drug: CC-401


```

```

import matplotlib.pyplot as plt
unique, counts = np.unique(y_train, return_counts=True)
drug_counts = dict(zip(label_encoder.inverse_transform(unique), counts))

plt.bar(drug_counts.keys(), drug_counts.values())
plt.xticks(rotation=90)
plt.title("Drug Distribution in Training Data")
plt.show()

```



This indicates bias towards certain drug, possibly CC-401.

```

import pandas as pd
import numpy as np
import tensorflow as tf
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.utils.class_weight import compute_class_weight
from tensorflow.keras.layers import Input, Dense, Flatten, Concatenate, Conv2D
from tensorflow.keras.models import Model

metadata_path = '/content/drive/My Drive/Senior Year/Applied Data Science/sample_reduced_features.csv'
metadata = pd.read_csv(metadata_path)

image_dir = '/content/drive/My Drive/Senior Year/Applied Data Science/Data/downsampled_data/'

metadata['Image_Path'] = metadata['Simplified_FileName'] + '_median_aggregated.tiff'
image_paths = image_dir + metadata['Image_Path']

meta_features = ['area', 'mean_intensity', 'bbox-0', 'bbox-1', 'bbox-2', 'bbox-3', 'eccentricity', 'solidity']
X_meta = metadata[meta_features]
y_labels = metadata['Metadata_pert_iname']

label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y_labels)

scaler = StandardScaler()
X_meta_scaled = scaler.fit_transform(X_meta)

X_meta_train, X_meta_val, y_train, y_val = train_test_split(X_meta_scaled, y_encoded, test_size=0.2, random_state=42)
image_paths_train, image_paths_val = train_test_split(image_paths, test_size=0.2, random_state=42)

class_weights = compute_class_weight('balanced', classes=np.unique(y_train), y=y_train)
class_weight_dict = dict(enumerate(class_weights))

def load_image(path):
    img = tf.keras.utils.load_img(path, target_size=(64, 64))
    return tf.keras.utils.img_to_array(img) / 255.0

def hybrid_data_generator(image_paths, X_meta, labels, class_weights, batch_size=32):
    num_samples = len(labels)
    while True:
        for offset in range(0, num_samples, batch_size):
            batch_paths = image_paths.iloc[offset:offset + batch_size].values
            batch_images = np.array([load_image(path) for path in batch_paths])
            batch_meta = X_meta[offset:offset + batch_size]
            batch_labels = labels[offset:offset + batch_size]
            batch_weights = np.array([class_weights[label] for label in batch_labels])
            yield {"image_input": batch_images, "meta_input": batch_meta}, batch_labels, batch_weights

image_input = Input(shape=(64, 64, 3), name="image_input")
x = Conv2D(32, (3, 3), activation='relu')(image_input)
x = Flatten()(x)

meta_input = Input(shape=(X_meta_train.shape[1],), name="meta_input")
meta_features = Dense(32, activation='relu')(meta_input)

combined = Concatenate()([x, meta_features])
output = Dense(len(label_encoder.classes_), activation='softmax')(combined)

hybrid_model = Model(inputs={"image_input": image_input, "meta_input": meta_input}, outputs=output)
hybrid_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

train_gen = hybrid_data_generator(image_paths_train, X_meta_train, y_train, class_weight_dict, batch_size=32)
val_data = {
    "image_input": np.array([load_image(path) for path in image_paths_val[:32]]),
    "meta_input": X_meta_val[:32]
}

history = hybrid_model.fit(
    train_gen,
    steps_per_epoch=len(y_train) // 32,
    validation_data=(val_data, y_val[:32]),
    epochs=30,
    callbacks=[
        tf.keras.callbacks.EarlyStopping(patience=5, restore_best_weights=True),
        tf.keras.callbacks.ReduceLROnPlateau(factor=0.2, patience=3, min_lr=1e-5)
    ],
    verbose=1

```

)

```

Epoch 1/30
362/362 ————— 182s 496ms/step - accuracy: 0.4018 - loss: 3.6175 - val_accuracy: 0.9375 - val_loss: 0.2972 - l
Epoch 2/30
362/362 ————— 180s 494ms/step - accuracy: 0.9855 - loss: 0.0831 - val_accuracy: 0.9688 - val_loss: 0.2776 - l
Epoch 3/30
362/362 ————— 179s 496ms/step - accuracy: 0.9988 - loss: 0.0125 - val_accuracy: 0.9688 - val_loss: 0.3002 - l
Epoch 4/30
362/362 ————— 179s 495ms/step - accuracy: 0.9998 - loss: 0.0102 - val_accuracy: 0.9688 - val_loss: 0.3027 - l
Epoch 5/30
362/362 ————— 179s 494ms/step - accuracy: 0.9998 - loss: 0.0099 - val_accuracy: 0.9688 - val_loss: 0.3020 - l
Epoch 6/30
362/362 ————— 179s 494ms/step - accuracy: 0.9998 - loss: 0.0098 - val_accuracy: 0.9688 - val_loss: 0.3038 - l
Epoch 7/30
362/362 ————— 181s 500ms/step - accuracy: 0.9998 - loss: 0.0097 - val_accuracy: 0.9688 - val_loss: 0.3055 - l

```

```

import os
import numpy as np
import tensorflow as tf
from tensorflow.keras.utils import load_img, img_to_array

example_data_dir = '/content/drive/My Drive/Senior Year/Applied Data Science/Data/example_data/'
example_image_paths = [os.path.join(example_data_dir, fname) for fname in os.listdir(example_data_dir) if fname.endswith('.tiff')]

def load_example_images(image_paths):
    images = [img_to_array(load_img(path, target_size=(64, 64))) / 255.0 for path in image_paths]
    return np.array(images)

example_images = load_example_images(example_image_paths)

predictions = hybrid_model.predict({"image_input": example_images, "meta_input": np.zeros((len(example_images), X_meta_train.shape[1]))})

predicted_labels = label_encoder.inverse_transform(np.argmax(predictions, axis=1))

for img_path, pred_label in zip(example_image_paths, predicted_labels):
    print(f"Image: {os.path.basename(img_path)}, Predicted Drug: {pred_label}")

```

```

1/1 ————— 0s 302ms/step
Image: r04c08f05p01-compound-FK866.tiff, Predicted Drug: SGX523
Image: r04c14f05p01-compound-DMS0.tiff, Predicted Drug: SGX523
Image: r06c10f05p01-compound-quinidine.tiff, Predicted Drug: SGX523
Image: r12c09f05p01-compound-FK866.tiff, Predicted Drug: SGX523
Image: r13c02f05p01-compound-LY2109761.tiff, Predicted Drug: SGX523

```

```

import pandas as pd
import os
import numpy as np
import tensorflow as tf
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from tensorflow.keras.layers import Input, Conv2D, Dense, Flatten, MaxPooling2D
from tensorflow.keras.models import Model

metadata_path = '/content/drive/My Drive/Senior Year/Applied Data Science/sample_reduced_features.csv'
metadata = pd.read_csv(metadata_path)

image_dir = '/content/drive/My Drive/Senior Year/Applied Data Science/Data/downsampled_data/'

metadata['Image_Path'] = image_dir + metadata['Simplified_FileName'] + '_median_aggregated.tiff'
image_paths = metadata['Image_Path']

y_labels = metadata['Metadata_pert_iname']

label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y_labels)

image_paths_train, image_paths_val, y_train, y_val = train_test_split(
    image_paths, y_encoded, test_size=0.2, random_state=42
)

def load_image(path):
    img = tf.keras.utils.load_img(path, target_size=(64, 64))
    return tf.keras.utils.img_to_array(img) / 255.0

def image_data_generator(image_paths, labels, batch_size=32):
    num_samples = len(labels)
    while True:
        for offset in range(0, num_samples, batch_size):
            batch_paths = image_paths.iloc[offset:offset + batch_size].values
            batch_images = np.array([load_image(path) for path in batch_paths])
            batch_labels = labels[offset:offset + batch_size]
            yield batch_images, batch_labels

image_input = Input(shape=(64, 64, 3), name="image_input")
x = Conv2D(16, (3, 3), activation='relu')(image_input)
x = MaxPooling2D(pool_size=(2, 2))(x)
x = Conv2D(32, (3, 3), activation='relu')(x)
x = MaxPooling2D(pool_size=(2, 2))(x)
x = Flatten()(x)
output = Dense(len(label_encoder.classes_), activation='softmax')(x)

simplified_model = Model(inputs=image_input, outputs=output)
simplified_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

train_gen = image_data_generator(image_paths_train, y_train, batch_size=32)
val_gen = image_data_generator(image_paths_val, y_val, batch_size=32)

history = simplified_model.fit(
    train_gen,
    steps_per_epoch=len(y_train) // 32,
    validation_data=val_gen,
    validation_steps=len(y_val) // 32,
    epochs=20,
    callbacks=[
        tf.keras.callbacks.EarlyStopping(patience=3, restore_best_weights=True),
        tf.keras.callbacks.ReduceLROnPlateau(factor=0.2, patience=2, min_lr=1e-5)
    ],
    verbose=1
)

```

Epoch 1/20
 41/362 ————— 2:35 484ms/step – accuracy: 0.1352 – loss: 5.3415

KeyboardInterrupt Traceback (most recent call last)

<ipython-input-26-c6744bbbe8ee> in <cell line: 66>()

64

65 # Train the model

----> 66 history = simplified_model.fit(

67 train_gen,

68 steps_per_epoch=len(y_train) // 32,

⏮ 10 frames

/usr/local/lib/python3.10/dist-packages/tensorflow/python/eager/execute.py in quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)

51 try:

52 ctx.ensure_initialized()

----> 53 tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_name,

54 inputs, attrs, num_outputs)

55 except core._NotOkStatusException as e:

KeyboardInterrupt:

import pandas as pd

metadata_path = '/content/drive/My Drive/Senior Year/Applied Data Science/sample_reduced_features.csv'

metadata = pd.read_csv(metadata_path)

drug_counts = metadata['Metadata_pert_iname'].value_counts()

top_5_drugs = drug_counts.head(5)

print("Top 5 most frequent drugs in the dataset:")

print(top_5_drugs)

import matplotlib.pyplot as plt

plt.figure(figsize=(8, 5))

top_5_drugs.plot(kind='bar', color='skyblue')

plt.title('Top 5 Most Frequent Drugs in the Dataset')

plt.xlabel('Drug')

plt.ylabel('Frequency')

plt.xticks(rotation=45)

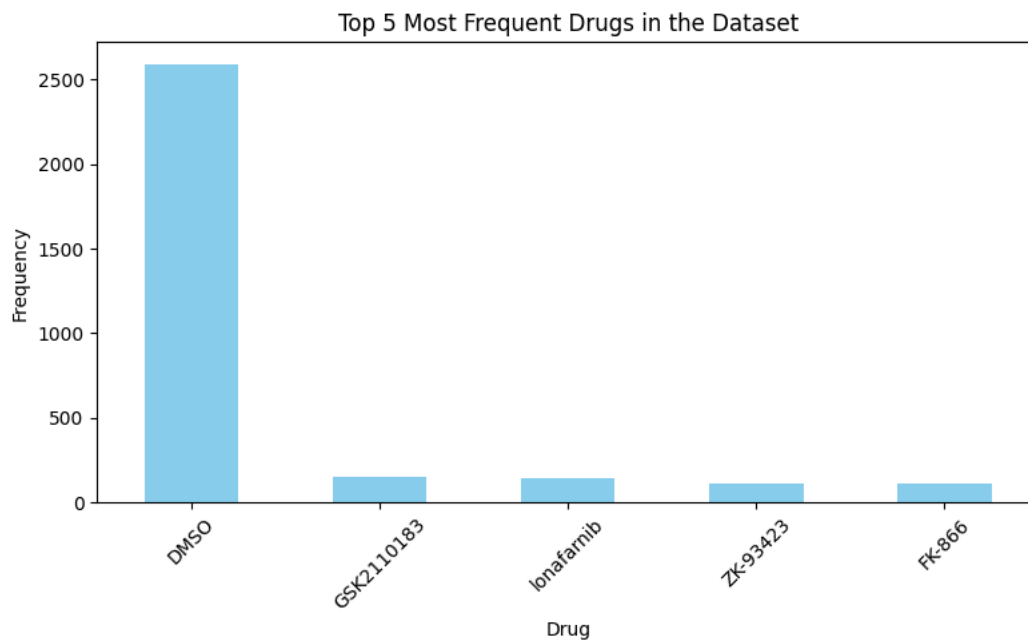
plt.tight_layout()

plt.show()

```

Top 5 most frequent drugs in the dataset:
Metadata_pert_iname
DMSO          2592
GSK2110183    153
lonafarnib    146
ZK-93423      113
FK-866        109
Name: count, dtype: int64

```



```
!pip install cellpose
```

```

from cellpose import models, io
import os
import pandas as pd
import numpy as np
from skimage.measure import regionprops, label
from tqdm import tqdm

\image_dir = '/content/drive/My Drive/Senior Year/Applied Data Science/Data/downsampled_data'
output_csv_path = '/content/drive/My Drive/Senior Year/Applied Data Science/cellpose_features.csv'

model = models.Cellpose(gpu=True, model_type='cyto')

features = []

for filename in tqdm(os.listdir(image_dir)):
    if filename.endswith('.tiff'):
        img_path = os.path.join(image_dir, filename)
        img = io.imread(img_path)

        outputs = model.eval(img, channels=[0, 0], diameter=None, flow_threshold=0.4, cellprob_threshold=0.0)
        masks = outputs[0]

        props = regionprops(label(masks), intensity_image=img)
        for prop in props:
            features.append({
                'Image': filename,
                'Area': prop.area,
                'Mean Intensity': prop.mean_intensity,
                'Bounding Box Sum': sum(prop.bbox),
                'Eccentricity': prop.eccentricity,
                'Solidity': prop.solidity
            })

df_features = pd.DataFrame(features)
df_features.to_csv(output_csv_path, index=False)

print(f"Feature extraction complete. Results saved to {output_csv_path}")

```



```

Requirement already satisfied: cellpose in /usr/local/lib/python3.10/dist-packages (3.1.0)
Requirement already satisfied: numpy<2.1,>=1.20.0 in /usr/local/lib/python3.10/dist-packages (from cellpose) (1.26.4)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from cellpose) (1.13.1)
Requirement already satisfied: natsort in /usr/local/lib/python3.10/dist-packages (from cellpose) (8.4.0)
Requirement already satisfied: tifffile in /usr/local/lib/python3.10/dist-packages (from cellpose) (2024.9.20)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from cellpose) (4.66.6)
Requirement already satisfied: numba>=0.53.0 in /usr/local/lib/python3.10/dist-packages (from cellpose) (0.60.0)
Requirement already satisfied: llvmlite in /usr/local/lib/python3.10/dist-packages (from cellpose) (0.43.0)
Requirement already satisfied: torch>=1.6 in /usr/local/lib/python3.10/dist-packages (from cellpose) (2.5.0+cu121)
Requirement already satisfied: opencv-python-headless in /usr/local/lib/python3.10/dist-packages (from cellpose) (4.10.0.84)
Requirement already satisfied: fastremap in /usr/local/lib/python3.10/dist-packages (from cellpose) (1.15.0)
Requirement already satisfied: imagecodecs in /usr/local/lib/python3.10/dist-packages (from cellpose) (2024.9.22)
Requirement already satisfied: roifile in /usr/local/lib/python3.10/dist-packages (from cellpose) (2024.9.15)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch>=1.6->cellpose) (3.16.1)
Requirement already satisfied: typing-extensions>=4.8.0 in /usr/local/lib/python3.10/dist-packages (from torch>=1.6->cellpos)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch>=1.6->cellpose) (3.4.2)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.10/dist-packages (from torch>=1.6->cellpose) (3.1.4)
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from torch>=1.6->cellpose) (2024.10.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.10/dist-packages (from torch>=1.6->cellpose) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from sympy==1.13.1->torch>=1.6)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from Jinja2->torch>=1.6->cellpose)
100% |██████████| 2867/2867 [1:24:09<00:00, 1.76s/it]
Feature extraction complete. Results saved to /content/drive/My Drive/Senior Year/Applied Data Science/cellpose_features.csv

```

```

import pandas as pd

cellpose_features_path = '/content/drive/My Drive/Senior Year/Applied Data Science/cellpose_features.csv'
metadata_path = '/content/drive/My Drive/Senior Year/Applied Data Science/Data/metadata_BR00116991.csv'

cellpose_features = pd.read_csv(cellpose_features_path)
metadata_features = pd.read_csv(metadata_path)

cellpose_features['Trimmed_Image'] = cellpose_features['Image'].str[:9]

metadata_features['Trimmed_FileName'] = metadata_features['FileName_OrigRNA'].str[:9]

merged_features = pd.merge(
    cellpose_features,
    metadata_features[['Trimmed_FileName', 'Metadata_pert_iname']],
    left_on='Trimmed_Image',
    right_on='Trimmed_FileName',
    how='left'
).drop(columns=['Trimmed_FileName', 'Trimmed_Image']) \

output_path = '/content/drive/My Drive/Senior Year/Applied Data Science/merged_features_metadata.csv'
merged_features.to_csv(output_path, index=False)

print(f"Merged dataset saved to {output_path}")
print(f"Final merged dataset has {len(merged_features)} rows.")

```

```

Merged dataset saved to /content/drive/My Drive/Senior Year/Applied Data Science/merged_features_metadata.csv
Final merged dataset has 396579 rows.

```

```

import os
import pandas as pd
import tensorflow as tf

merged_features_path = '/content/drive/My Drive/Senior Year/Applied Data Science/merged_features_metadata.csv'
image_dir = '/content/drive/My Drive/Senior Year/Applied Data Science/Data/downsampled_data/'

merged_features = pd.read_csv(merged_features_path)

merged_features['Image'] = merged_features['Image'].apply(lambda x: x if x.endswith('.tiff') else x + '.tiff')

merged_features['Image'] = merged_features['Image'].apply(lambda x: os.path.join(image_dir, x))

merged_features = merged_features[merged_features['Image'].apply(os.path.exists)]

print(f"Number of images matched: {merged_features.shape[0]}")

```

```

Number of images matched: 396579

```

```

import os
import pandas as pd
import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler

merged_features_path = '/content/drive/My Drive/Senior Year/Applied Data Science/merged_features_metadata.csv'
merged_features = pd.read_csv(merged_features_path)

aggregated_features = merged_features.groupby('Image').agg({
    'Area': 'mean',
    'Mean Intensity': 'mean',
    'Bounding Box Sum': 'mean',
    'Eccentricity': 'mean',
    'Solidity': 'mean',
    'Metadata_pert_iname': 'first'
}).reset_index()

label_encoder = LabelEncoder()
aggregated_features['Encoded_Label'] = label_encoder.fit_transform(aggregated_features['Metadata_pert_iname'])

X = aggregated_features[['Area', 'Mean Intensity', 'Bounding Box Sum', 'Eccentricity', 'Solidity']].values
y = aggregated_features['Encoded_Label'].values

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_val, y_train, y_val = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

model = tf.keras.Sequential([
    tf.keras.layers.Dense(64, activation='relu', input_shape=(X_train.shape[1],)),
    tf.keras.layers.Dense(32, activation='relu'),
    tf.keras.layers.Dense(len(label_encoder.classes_), activation='softmax')
])

model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

history = model.fit(
    X_train, y_train,
    validation_data=(X_val, y_val),
    epochs=30,
    batch_size=32,
    callbacks=[
        tf.keras.callbacks.EarlyStopping(patience=5, restore_best_weights=True),
        tf.keras.callbacks.ReduceLROnPlateau(factor=0.2, patience=3, min_lr=1e-5)
    ],
    verbose=1
)

model.save('/content/drive/My Drive/Senior Year/Applied Data Science/aggregated_model.h5')

print("Training complete and model saved.")

```

```

Epoch 1/30
/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input_shape` to `Input` layer.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
72/72 _____ 3s 30ms/step - accuracy: 0.0308 - loss: 5.5161 - val_accuracy: 0.1690 - val_loss: 5.3892 - learnin
Epoch 2/30
72/72 _____ 0s 2ms/step - accuracy: 0.1802 - loss: 5.2335 - val_accuracy: 0.1690 - val_loss: 5.0211 - learnin
Epoch 3/30
72/72 _____ 0s 2ms/step - accuracy: 0.1710 - loss: 4.8335 - val_accuracy: 0.1725 - val_loss: 4.8830 - learnin
Epoch 4/30
72/72 _____ 0s 2ms/step - accuracy: 0.1773 - loss: 4.6448 - val_accuracy: 0.1777 - val_loss: 4.8035 - learnin
Epoch 5/30
72/72 _____ 0s 2ms/step - accuracy: 0.1920 - loss: 4.4562 - val_accuracy: 0.1812 - val_loss: 4.7424 - learnin
Epoch 6/30
72/72 _____ 0s 2ms/step - accuracy: 0.2061 - loss: 4.3107 - val_accuracy: 0.1847 - val_loss: 4.6666 - learnin
Epoch 7/30
72/72 _____ 0s 2ms/step - accuracy: 0.2088 - loss: 4.2251 - val_accuracy: 0.1882 - val_loss: 4.6080 - learnin
Epoch 8/30
72/72 _____ 0s 2ms/step - accuracy: 0.1926 - loss: 4.2008 - val_accuracy: 0.1847 - val_loss: 4.5728 - learnin
Epoch 9/30
72/72 _____ 0s 2ms/step - accuracy: 0.2144 - loss: 4.0175 - val_accuracy: 0.1829 - val_loss: 4.5603 - learnin
Epoch 10/30
72/72 _____ 0s 2ms/step - accuracy: 0.2281 - loss: 3.9656 - val_accuracy: 0.1899 - val_loss: 4.5288 - learnin
Epoch 11/30
72/72 _____ 0s 2ms/step - accuracy: 0.2243 - loss: 3.9069 - val_accuracy: 0.1847 - val_loss: 4.5091 - learnin
Epoch 12/30

```

```

72/72 ————— 0s 2ms/step - accuracy: 0.2155 - loss: 3.8868 - val_accuracy: 0.1899 - val_loss: 4.5241 - learnin
Epoch 13/30
72/72 ————— 0s 2ms/step - accuracy: 0.2197 - loss: 3.8816 - val_accuracy: 0.1899 - val_loss: 4.5174 - learnin
Epoch 14/30
72/72 ————— 0s 2ms/step - accuracy: 0.2266 - loss: 3.8524 - val_accuracy: 0.1934 - val_loss: 4.5183 - learnin
Epoch 15/30
72/72 ————— 0s 2ms/step - accuracy: 0.2346 - loss: 3.7359 - val_accuracy: 0.1969 - val_loss: 4.4996 - learnin
Epoch 16/30
72/72 ————— 0s 2ms/step - accuracy: 0.2412 - loss: 3.7286 - val_accuracy: 0.1934 - val_loss: 4.5026 - learnin
Epoch 17/30
72/72 ————— 0s 2ms/step - accuracy: 0.2374 - loss: 3.7498 - val_accuracy: 0.1951 - val_loss: 4.5040 - learnin
Epoch 18/30
72/72 ————— 0s 2ms/step - accuracy: 0.2372 - loss: 3.7501 - val_accuracy: 0.1951 - val_loss: 4.5038 - learnin
Epoch 19/30
72/72 ————— 0s 2ms/step - accuracy: 0.2458 - loss: 3.6578 - val_accuracy: 0.1951 - val_loss: 4.5054 - learnin
Epoch 20/30
72/72 ————— 0s 2ms/step - accuracy: 0.2293 - loss: 3.7287 - val_accuracy: 0.1951 - val_loss: 4.5055 - learnin
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file for
Training complete and model saved.

```

```
import matplotlib.pyplot as plt
```

```



plt.figure(figsize=(10, 5))
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy', linestyle='--')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig('/content/drive/My Drive/Senior Year/Applied Data Science/training_validation_accuracy.png')
plt.show()

```

```

plt.figure(figsize=(10, 5))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss', linestyle='--')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig('/content/drive/My Drive/Senior Year/Applied Data Science/training_validation_loss.png')
plt.show()

```



```
import os
```

```
from google.colab import drive
```